

Deep Learning Kernel Based Model Order Reduction of Millimeter-Wave Passive Devices

Nathan T. Tilau

Department of Electrical and Computer Engineering

University Name

City, Country

Email: ntilau@example.com

Abstract—Full-wave electromagnetic simulation of millimeter-wave passive devices via the finite element method (FEM) is computationally expensive, particularly for parametric sweeps over material and frequency dimensions required during design optimization. This paper presents a surrogate modeling framework based on Deep Kernel Learning (DKL) that accelerates S-parameter prediction for a bilateral finline filter by over three orders of magnitude while maintaining less than 0.25% relative error. The approach trains independent DNN-to-GP mappings—Deep Kernel Learning following Wilson et al. 2016—for each complex scattering parameter, employing a product of RBF kernels over learned deep features to capture the wave physics. An eight-dimensional harmonic input encoding enables frequency-aware feature extraction without explicit wavenumber computation. We describe the complete pipeline from FEM data generation through model training, validation, and uncertainty quantification, achieving 0.19% relative error on S_{11} and 0.21% on S_{21} using only 4050 training samples from 50 dielectric constant values across 81 frequency points.

Index Terms—Model order reduction, deep kernel learning, Gaussian process, finite element method, waveguide filter, surrogate modeling, millimeter-wave

I. INTRODUCTION

Millimeter-wave passive components—filters, couplers, circulators, and antennas—require accurate electromagnetic characterization across wide frequency bands and varying material parameters. The finite element method (FEM) is a workhorse for such analysis, but its computational cost scales poorly with electrical size and the number of parametric evaluations needed for design optimization, tolerance analysis, and yield estimation.

Model order reduction (MOR) addresses this bottleneck by constructing lightweight surrogates that approximate the input–output map of the full-wave solver. Traditional approaches include proper orthogonal decomposition (POD) [1], reduced-basis methods [2], and rational interpolation (Vector Fitting) [3]. More recently, machine learning surrogates—particularly neural networks [4], [5] and Gaussian processes [6]—have gained traction for their ability to handle nonlinear parameter dependencies.

Deep Kernel Learning (DKL) [7] combines the representational power of deep neural networks with the probabilistic

calibration of Gaussian processes. A DNN learns a low-dimensional feature embedding from the raw inputs, and a GP with a base kernel operates on those learned features. This hybrid retains the GP’s uncertainty quantification and data efficiency while leveraging the DNN’s ability to discover relevant structure.

This paper presents a DKL-based surrogate for the bilateral finline filter, a canonical millimeter-wave component with significant parametric sensitivity. Our contributions are:

- 1) A complete, end-to-end MOR pipeline from FEM data generation to trained surrogate with uncertainty estimates.
- 2) An input encoding strategy using harmonic features (sin/cos expansions of frequency) that captures the wave-like propagation physics.
- 3) A targeted output transformation: direct Re/Im regression for S_{11} (reflection), and log-magnitude plus detrended unwrapped-phase regression for S_{21} (transmission), which handles a $14\times$ amplitude variation between passband and notch.
- 4) Demonstration of sub-0.25% relative error on both S-parameters using a product-of-RBF-kernels DKL architecture.

II. PROBLEM FORMULATION

A. Parameterized Waveguide Scattering

Consider a bilateral finline filter [8] consisting of a rectangular waveguide loaded with a dielectric substrate patterned by a metallic finline. The geometry is defined by a substrate of relative permittivity ϵ_r and height h , excited by the TE_{10} mode at frequency f . For each input pair $\mathbf{x} = (\epsilon_r, f)$, the FEM solver [9] assembles the discrete Helmholtz system

$$(\mathbf{S} - k^2 \mathbf{T}) \mathbf{u} = \mathbf{b}, \quad (1)$$

where $\mathbf{S}$ and $\mathbf{T}$ are the stiffness and mass matrices, $k = 2\pi f/c_0$ is the free-space wavenumber, $\mathbf{u}$ is the unknown electric field, and $\mathbf{b}$ encodes waveguide port excitation. After solving (1), the modal expansion at each port yields the complex scattering parameters $S_{11}(\mathbf{x})$ and $S_{21}(\mathbf{x})$.

For a parametric sweep over $N_{\varepsilon_r} = 50$ permittivity values and $N_f = 81$ frequency points (4050 solves total), each solve requires $\mathcal{O}(N_{\text{dof}}^2)$ operations where $N_{\text{dof}} \sim 10^4$ for a second-order hierarchical basis. The total wall-clock time is on the order of hours on a desktop workstation—prohibitively expensive for the hundreds of thousands of evaluations needed in optimization or Monte Carlo tolerance analysis.

B. Goal: Data-Efficient Surrogate

We seek a surrogate f_θ such that for any (ε_r, f) in the design space, $S_{11} \approx f_\theta^{(1)}(\varepsilon_r, f)$ and $S_{21} \approx f_\theta^{(2)}(\varepsilon_r, f)$, evaluated in microseconds rather than seconds, with predictive uncertainty quantification and without a separate mesh or solver invocation.

III. DEEP KERNEL LEARNING SURROGATE

A. Architecture

Deep Kernel Learning [7] constructs a kernel $k(\mathbf{x}, \mathbf{x}')$ that is itself a learned function of the inputs via a deep neural network. Given a base kernel $k_{\text{base}}(\cdot, \cdot)$ (e.g., RBF, Matern), the DKL kernel is

$$k_\theta(\mathbf{x}, \mathbf{x}') = k_{\text{base}}(g_\phi(\mathbf{x}), g_\phi(\mathbf{x}')), \quad (2)$$

where $g_\phi : \mathbb{R}^d \rightarrow \mathbb{R}^p$ is a DNN parameterized by ϕ , and $\theta = (\phi, \psi)$ with ψ the base kernel hyperparameters. The model is an ExactGP [10] with constant mean function $\mu(\mathbf{x}) = \mu_0$ and Gaussian likelihood $y(\mathbf{x}) \sim \mathcal{GP}(\mu(\mathbf{x}), k_\theta(\mathbf{x}, \mathbf{x}') + \sigma_n^2 \delta_{\mathbf{x}\mathbf{x}'}).$

For each S-parameter channel we train an independent DNN-GP model. The DNN feature extractor architecture is

$$\mathbb{R}^8 \xrightarrow{\text{Linear}} \mathbb{R}^{1000} \xrightarrow{\text{ReLU}} \mathbb{R}^{1000} \xrightarrow{\text{ReLU}} \mathbb{R}^{500} \xrightarrow{\text{ReLU}} \mathbb{R}^{64}. \quad (3)$$

The base kernel is a *product* of two multi-dimensional RBF kernels with Automatic Relevance Determination (ARD):

$$k_{\text{base}}(\mathbf{z}, \mathbf{z}') = \sigma_f^2 \exp \left(-\frac{1}{2} \sum_{j=1}^p \frac{(z_j - z'_j)^2}{\ell_{1,j}^2} \right) \times \exp \left(-\frac{1}{2} \sum_{j=1}^p \frac{(z_j - z'_j)^2}{\ell_{2,j}^2} \right). \quad (4)$$

The product of two RBF kernels captures interactions at different length scales in the learned feature space—analogueous to a mixture-of-kernels approach that can represent both smooth and rapidly-varying responses.

B. Input Encoding

The raw inputs (ε_r, f) are not well suited for direct regression; the wave physics is oscillatory in f with no simple linear relationship. We construct an eight-dimensional feature vector:

$$\mathbf{x} = [\varepsilon_r, f/f_0, \sin \omega, \cos \omega, \sin 2\omega, \cos 2\omega, \sin 3\omega, \cos 3\omega]^\top, \quad (5)$$

where $\omega = 2\pi f/f_0$ with $f_0 = 10$ GHz reduces the raw frequency to a moderate angular range. The sin/cos pairs of the first three harmonics provide the GP with explicit basis functions for sinusoidal variation—critical for modeling standing-wave patterns in the filter’s passband without requiring the DNN to discover Fourier features from scratch. The scaled frequency f/f_0 retains the monotonic trend. All eight features are z-score normalized before training.

C. Output Transforms

The complex S-parameters have fundamentally different characteristics: S_{11} (reflection) undergoes a deep null at the filter’s passband where its amplitude drops below -45 dB, while S_{21} (transmission) spans a $14\times$ amplitude range. Despite these differences, we adopt a *unified* Log-Magnitude + Detrended Phase representation for both parameters. For S_{11} , the phase model is trained only on points where $|S_{11}| > 0.05$ to avoid the numerically ill-defined region near the reflection null.

1) *Log-Magnitude + Detrended Phase (Both S_{11} and S_{21}):* Each complex S-parameter is encoded as two real-valued targets:

$$y_1^{(p)} = \log_{10} |S_p|, \quad (6)$$

$$y_2^{(p)} = \text{unwrap}[\angle S_p] - (\alpha_p f + \beta_p), \quad (7)$$

where $p \in \{11, 21\}$ indexes the S-parameter and the linear phase trend $\alpha_p f + \beta_p$ is fitted per ε_r value via least squares [11] and stored for reconstruction. This *detrended* unwrapped phase removes the smooth group delay, isolating the frequency-dependent fine structure that carries information about resonance and coupling.

During inference, the log-magnitude is exponentiated and the phase is restored:

$$S_p^{(\text{pred})} = 10^{y_1^{(p)}} \exp[j(y_2^{(p)} + \alpha_p f + \beta_p)]. \quad (8)$$

D. Training Procedure

Training proceeds in two stages per model, following the DKL protocol [7]:

1) *Stage 1: DNN Warm-Start:* The feature extractor is pretrained via MSE loss to predict the target using the mean of its 64-dimensional output. This stage (up to 100 epochs, AdamW optimizer, learning rate 10^{-3}) initializes the DNN weights in a region where the GP can immediately exploit meaningful features.

2) *Stage 2: Joint Marginal Likelihood Optimization:* The full model (DNN + GP kernel + likelihood) is trained jointly by minimizing the negative log marginal likelihood:

$$\mathcal{L}(\theta) = -\log p(\mathbf{y} | \mathbf{X}, \theta) = \frac{1}{2} \mathbf{y}^\top \mathbf{K}_\theta^{-1} \mathbf{y} + \frac{1}{2} \log |\mathbf{K}_\theta| + \frac{N}{2} \log 2\pi, \quad (9)$$

where $\mathbf{K}_\theta = k_\theta(\mathbf{X}, \mathbf{X}) + \sigma_n^2 \mathbf{I}$ and N is the number of training points. Optimizer: AdamW with learning rate 5×10^{-3} , weight decay 10^{-4} , and cosine annealing scheduler. Training runs for 500 epochs.




Fig. 1. DKL surrogate reflection magnitude $|S_{11}|$ (dB). Solid: FEM ground truth; dashed: DKL prediction. Bold traces: held-out ϵ_r values. Color maps $\epsilon_r \in [2.0, 2.2]$.




Fig. 2. DKL surrogate reflection phase $\angle S_{11}$ (degrees). Solid: FEM; dashed: DKL. Only points with both $|S_{11}| > 0.05$ are shown to suppress noise near the reflection null at ~ 150 GHz.

E. Validation Strategy

The dataset comprises $N_{\epsilon_r} \times N_f$ samples. To avoid overly optimistic interpolation metrics, validation is performed on *held-out* ϵ_r values rather than random splits. At most one-third of ϵ_r values (capped at 10) are held out; the model never sees any frequency point for those ϵ_r values during training. This tests the surrogate’s ability to extrapolate across material parameters—the most practically relevant generalization task.

F. Uncertainty Quantification

As a Bayesian model, the trained GP produces a full predictive distribution:

$$p(y^* | \mathbf{x}^*, \mathcal{D}) = \mathcal{N}(\mu_*, k_*^\top \mathbf{K}_\theta^{-1} k_* + k_{**} - k_*^\top \mathbf{K}_\theta^{-1} k_*), \quad (10)$$

where $k_* = k_\theta(\mathbf{X}, \mathbf{x}^*)$ and $k_{**} = k_\theta(\mathbf{x}^*, \mathbf{x}^*)$. The predictive standard deviation at $\mathbf{x}^*$ provides a per-point confidence interval at negligible additional cost.

IV. RESULTS AND DISCUSSION

A. Data Generation

Training data were generated using the pyFES finite element solver [12] with a second-order hierarchical basis (pOrd = 2, hOrd = 1, approximately 13,000 DOFs). The bilateral filter geometry consists of a WR-6 rectangular waveguide ($a = 1.651$ mm, $b = 0.8255$ mm) housing a dielectric substrate ($\epsilon_r \in [2.0, 2.2]$, $h = 0.8255$ mm) with a unilateral finline. Fifteen waveguide modes were retained per port. The frequency sweep spans 138–158 GHz in 81 equidistant steps, covering both the TE_{10} passband and the filter’s rejection notch.

TABLE I
VALIDATION METRICS FOR DKL SURROGATE.

Parameter	RMSE	Rel. error (%)	Phase error ($^\circ$)
S_{11}	1.0×10^{-3}	0.19	0.14
S_{21}	1.7×10^{-3}	0.21	0.05

B. Prediction Accuracy

Table I summarizes validation metrics averaged over held-out ϵ_r values.

Both S-parameters achieve relative errors well below 1%. The reflection coefficient S_{11} benefits from the direct Re/Im representation, which avoids the nonlinearities inherent in magnitude-phase conversion. The transmission S_{21} error remains low despite the $14\times$ amplitude range, validating the log-magnitude + detrended-phase strategy.

A single FEM solve requires approximately 2–3 seconds on an Apple M2 workstation. Generating the full 4050-sample training set takes approximately 2.5 hours. After training, the DKL surrogate evaluates 4050 points in under 100 milliseconds—a speedup factor exceeding 10^4 over the FEM solver. Training all four models (including data loading and preprocessing) completes in approximately 30 minutes on a single GPU (NVIDIA RTX 4070).

C. Effect of Input Encoding

To study the contribution of harmonic features, we trained an ablation variant using only $(\epsilon_r, f/f_0)$ as inputs. The full 8-dimensional encoding reduces S_{21} relative error from 0.61% to 0.23%. The sin/cos features are particularly important in the filter’s transition band where the derivative dS/df is largest; without them, the GP posterior underfits the rapid phase variation.




Fig. 3. S_{11} relative prediction error versus frequency and ϵ_r . Color scale is $\log_{10}(\text{error } \%)$.




Fig. 4. DKL surrogate transmission magnitude $|S_{21}|$ (dB). The $14\times$ amplitude range between passband (~ 0 dB) and notch (~ -20 dB) motivates the $\log_{10}|S_{21}|$ representation.

D. Uncertainty Calibration

Figures 7 and 8 illustrate the predictive uncertainty for S_{21} at two held-out ϵ_r values. The 95% confidence interval ($\pm 2\sigma$) correctly widens in frequency regions where the surrogate is less certain, reflecting the model's epistemic uncertainty.

V. PIPELINE OVERVIEW

Algorithm 1 summarizes the complete MOR workflow, from FEM data generation to deployment of the trained surrogate with uncertainty quantification.

VI. CONCLUSION

We have presented a Deep Kernel Learning framework for model order reduction of millimeter-wave passive devices, applied to the bilateral finline filter. We have presented a Deep Kernel Learning framework for model order reduction of millimeter-wave passive devices, applied to the bilateral finline




Fig. 5. DKL surrogate transmission phase $\angle S_{21}$ (degrees). Linear group delay is removed during training via per- ϵ_r detrending and restored during inference.




Fig. 6. S_{21} relative prediction error versus frequency and ϵ_r . The DKL model maintains sub-1% error across the full design space despite the large dynamic range.

filter. The DKL surrogate achieves sub-0.25% relative error on both reflection and transmission S-parameters while providing Bayesian uncertainty quantification and a $10^4\times$ speedup over the full-wave FEM solver. Key technical elements include: (1) harmonic input encoding that captures wave physics, (2) a product-of-RBF base kernel that models multi-scale feature interactions, (3) targeted output representations (direct Re/Im for S_{11} , log-magnitude+detrended-phase for S_{21}), and (4) a two-stage training protocol that warm-starts the DNN features before joint GP optimization.

The approach is general: the same pipeline applies to any parameterized FEM system with smooth input–output maps, including waveguide filters, couplers, and antenna arrays. Future work includes extending the surrogate to support multi-port devices (S_N -parameters) with a shared feature extractor, incorporating frequency-rational basis functions through physics-informed kernels, and exploring active learning strate-




Fig. 7. S_{21} predictive mean and $\pm 2\sigma$ confidence interval for $\varepsilon_r = 2.053$ (held-out). Shaded: $\pm 2\sigma$; dots: FEM ground truth.




Fig. 8. S_{21} predictive mean and $\pm 2\sigma$ confidence interval for $\varepsilon_r = 2.196$ (held-out). Shaded: $\pm 2\sigma$; dots: FEM ground truth.

gies that query the FEM solver only at points of high predictive uncertainty.

REFERENCES

- [1] A. Quarteroni, A. Manzoni, and F. Negri, *Reduced Basis Methods for Partial Differential Equations*. Springer, 2015.
- [2] G. Rozza, D. B. P. Huynh, and A. T. Patera, “Reduced basis approximation and a posteriori error estimation for affinely parametrized elliptic coercive partial differential equations,” *Arch. Comput. Methods Eng.*, vol. 15, no. 3, pp. 229–275, 2008.
- [3] B. Gustavsen and A. Semlyen, “Rational approximation of frequency domain responses by vector fitting,” *IEEE Trans. Power Del.*, vol. 14, no. 3, pp. 1052–1061, 1999.
- [4] P. Burrascano, S. Fiori, and M. Mongiardo, “A review of artificial neural networks for waveguide filter optimization,” *Int. J. RF Microw. Comput.-Aided Eng.*, vol. 8, no. 5, pp. 407–421, 1998.
- [5] F. Feng, C. Zhang, J. Ma, and Q.-J. Zhang, “Parametric modeling of microwave components using adjoint neural networks and pole-residue transfer functions with EM sensitivity analysis,” *IEEE Trans. Microw. Theory Tech.*, vol. 66, no. 4, pp. 1755–1768, 2018.
- [6] L. P. Swiler and T. L. Paez, “Gaussian process emulators for simulating complex physics models,” Sandia National Laboratories, Tech. Rep. SAND2011-2693, 2011.

Algorithm 1 Deep Kernel Learning MOR Pipeline

Require: Permittivity range $[\varepsilon_{\min}, \varepsilon_{\max}]$, frequency range $[f_{\min}, f_{\max}]$, mesh data for bilateral filter

Ensure: Trained DKL surrogate with uncertainty estimates

— Phase A: Data Generation —

- 1: **for** each ε_r value **do**
- 2: Read mesh; assign substrate ε_r ; set port BCs
- 3: Assemble $\mathbf{S}$, $\mathbf{T}$ matrices via ASSEMBLELINEAR
- 4: **for** each frequency f **do**
- 5: Assemble waveguide port matrices
- 6: Solve $(\mathbf{S} - k^2\mathbf{T})\mathbf{u} = \mathbf{b}$
- 7: Extract $\mathbf{S}_{11}$, $\mathbf{S}_{21}$ via modal expansion
- 8: Store $(\varepsilon_r, f, \mathbf{S}_{11}, \mathbf{S}_{21})$
- 9: **end for**
- 10: **end for**

— Phase B: Feature Construction —

- 11: Compute harmonic encoding: $\mathbf{x} \leftarrow [\varepsilon_r, f/f_0, \sin \omega, \cos \omega, \sin 2\omega, \cos 2\omega, \sin 3\omega, \cos 3\omega]^\top$
- 12: $y_1 \leftarrow \text{Re}(\mathbf{S}_{11})$, $y_2 \leftarrow \text{Im}(\mathbf{S}_{11})$
- 13: $y_3 \leftarrow \log_{10} |\mathbf{S}_{21}|$
- 14: Fit $\mathbf{S}_{21}$ phase slope/intercept per ε_r ; compute detrended y_4
- 15: z-score normalize $\mathbf{x}$, $\{y_i\}$ independently

— Phase C: Training (per channel $i = 1 \dots 4$) —

- 16: Initialize FeatureExtractor DNN ($8 \rightarrow 1000 \rightarrow 1000 \rightarrow 500 \rightarrow 64$)
- 17: Initialize ExactGP with product-RBF kernel + Gaussian-Likelihood

Stage 1 – DNN warm-start:

- 18: Optimize ϕ via MSE loss (100 epochs, $\text{lr} = 10^{-3}$)

Stage 2 – Joint DKL:

- 19: Optimize θ via $-\log p(\mathbf{y} | \mathbf{X}, \theta)$ (500 epochs, $\text{lr} = 5 \times 10^{-3}$, CosineAnnealing)
- 20: Evaluate on held-out ε_r values; store checkpoint

— Phase D: Inference —

- 21: **for** each query (ε_r, f) **do**
- 22: Compute harmonic features $\mathbf{x}$; z-score normalize
- 23: **for** each channel i **do**
- 24: $\mu_i, \sigma_i \leftarrow \text{ExactGP.predict}(\mathbf{x})$
- 25: $\hat{y}_i \leftarrow \text{inverse-normalize}(\mu_i)$
- 26: **end for**
- 27: $\hat{\mathbf{S}}_{11} \leftarrow \hat{y}_1 + j\hat{y}_2$
- 28: $|\hat{\mathbf{S}}_{21}| \leftarrow 10^{\hat{y}_3}$; restore phase from detrended $\hat{y}_4$
- 29: $\hat{\mathbf{S}}_{21} \leftarrow |\hat{\mathbf{S}}_{21}| e^{j\hat{y}_4}$
- 30: **end for**
- 31: **return** $\hat{\mathbf{S}}_{11}$, $\hat{\mathbf{S}}_{21}$, $\sigma_{\mathbf{S}_{11}}$, $\sigma_{\mathbf{S}_{21}}$

- [7] A. G. Wilson, Z. Hu, R. Salakhutdinov, and E. P. Xing, “Deep kernel learning,” in *Proc. 19th Int. Conf. Artificial Intelligence and Statistics (AISTATS)*, 2016, pp. 370–378.
- [8] J. Bornemann, U. Rosenberg, and S. Amari, *Planar Microwave Engineering: A Practical Guide to Theory, Measurement, and Circuits*. Cambridge University Press, 2010.
- [9] J.-M. Jin, *The Finite Element Method in Electromagnetics*, 3rd ed. Wiley, 2015.
- [10] C. E. Rasmussen and C. K. I. Williams, *Gaussian Processes for Machine Learning*. MIT Press, 2006.

- [11] S. M. Kay, *Fundamentals of Statistical Signal Processing: Estimation Theory*. Prentice Hall, 1998.
- [12] pyFES: Python Finite Element Solver for electromagnetics, 2025. GitHub repository.